

If Brand 30 costs as much to get right as Brand 1, you don't have a house of brands. **You have thirty startups sharing an office.**

Brand Zero — a shared memory that tells a launching brand what its neighbours already learned, and says plainly when they are no help.

Satyam Singh · 15 August 2026



The 31st brand.
It ships no product.

Think9's twelfth brand launches next month. Nobody knows what will work.

It sits in **ingestible wellness** — the category The Good Bug, SuperYou and Panchamrit are already in. The team must pick the message for its launch ads and put budget behind it. They believe it will lift sales about 35%.

WHAT THE LIFT COULD REALLY BE

−24% to +140%

A campaign that loses money, or one that nearly doubles sales. Both are consistent with what they know today.

COST OF GETTING THIS ONE CALL WRONG

₹5.01L

On ₹40L riding on the decision. Not a forecast of revenue — the expected cost of calling it the wrong way.

AND THIS REPEATS

× 19

Eleven brands live, thirty planned. Nineteen more launches, each paying close to full price for the same lesson.

The bottleneck is not "knowledge is fragmented." It is **launch decision latency** — the weeks and rupees between "*we're launching X*" and "*we know what works*" — paid again for every brand. Today that learning travels through people in meetings: slow, lossy, and it stops working around brand seven.

The obvious fix is worse than building nothing.

Pool all eleven brands, take the average, serve one number. This is the shape most "central AI brain" proposals quietly take. It does not make the answer more accurate — the estimate is still just the portfolio average. It only makes the answer **sound** certain. And certainty is the part that makes people spend.

WHAT IT CLAIMS

±11 pts

A tight, confident, specific planning range

HOW FAR OFF IT TYPICALLY IS

±164 pts

About 14× wider than it admits

HOW OFTEN IT IS ACTUALLY RIGHT

7%

It tells you 80%. Truth lands inside about one launch in fourteen

BRAND ZERO, SAME QUESTION


**−5% to
+91%**

Wider, and honest about it

Why this matters more than any win we could show. A wide honest answer sends the team to go and test. A narrow confident wrong one sends them to spend. Comics and probiotics do not share a price response — and a system that averages them together will never tell you so. Knowing when *not* to pool is not a feature of this system; it is the reason it exists.

Brand Zero — four organs and one loop

Nothing here asks anyone to change how they work. No one is asked to "log their decisions" — they won't. The system observes the decisions people were already making, from where they already make them.



All four organs are implemented and tested — github.com/SatyamSingh-Git/brand-zero · python demo.py runs the whole loop in ten seconds.

A claim, not a document. And not all claims count the same.

Two decisions carry most of the weight, and both are the kind of thing a system gets wrong silently.

- **Everything lands on one scale.** A price elasticity in face wash and one in a ₹4,000 device are not the same quantity. Effects arrive in whatever units the source used — a 1.19 conversion ratio, a 0.07 fractional change — and are converted to one common axis *before* anything is pooled. Skip this and the model does confident arithmetic on incompatible numbers.
- **Weak evidence still counts — it just counts less.** A randomised holdout and a line in a pricing meeting are both evidence. Each tier discounts how much a claim can move the answer. Nothing is thrown away; nothing weak can move a posterior on its own.
- **Every number traces back.** A claim without a link to its source is rejected by the schema. Any answer the system gives can be followed to the Slack message or spreadsheet row it came from, in two clicks.

WHERE IT CAME FROM	COUNTS AS
Randomised holdout / clean A-B	100%
Geo split, time split	70%
Before / after, no control	35%
Correlational read	15%
Said in a meeting	5%

The same claim, two ways. A pricing-meeting opinion with a nominal error of 0.40 enters the model at **1.79**. It is kept, it is auditable, and it cannot move the answer by itself.

It works out which brands are alike. Nobody tells it.

Three levels — **portfolio** → **category** → **brand** — fit separately for each kind of decision. Two levels would assume the portfolio is homogeneous, which is the assumption the whole system exists to test.

DOES WHAT ONE BRAND LEARNED MOVE BELIEF ABOUT ANOTHER?	AD MESSAGE	PRICE
The Good Bug → SuperYou	0.17	0.04
Amar Chitra Katha → Tinkle	0.16	0.01
Neude → Beauty by Bie	0.18	0.02
Tinkle → Neude	0.01	0.03
The Good Bug → Broadway	-0.07	-0.01

Shaded rows are brands in the same category. **Nobody supplied these numbers** — they fall out of how alike the brands turned out to be once the evidence was pooled. That is why Tinkle → Neude is allowed to be zero.

AD MESSAGES, WITHIN A CATEGORY

Pools

Neighbours remove 58% of the uncertainty — enough to plan against.

PRICE LADDERS, SAME BRANDS


Declines

Neighbours remove 6%. It returns a wide range and tells the team to run a real price test.

The ceiling, stated plainly. No amount of neighbour evidence pulls a new brand below how much brands genuinely differ. Transfer buys the distance from wide to narrow, and not one point further. A system claiming otherwise would be lying.

We built it, measured it, and it contradicted our own pitch.

The obvious chart for this deck is cost-per-correct-decision falling as brand number climbs. **The model does not produce that curve.** Portfolio size barely matters. What matters is how many *similar* brands a launching brand has.



How much a launching brand still doesn't know. Lower is better.

FIRST TWO SIMILAR BRANDS

–54%

uncertainty removed

GROWING 11 → 30 BRANDS

–12%

uncertainty removed

A house of brands compounds by going **deeper into fewer categories**, not wider across more.

Five categories of six brands learns faster than thirty categories of one. Thirty brands spread across thirty categories is, for learning purposes, thirty startups sharing an office after all.

`cold_start.py` — one world per seed, grown progressively, target brand never observed, averaged over six worlds. The first version of this measurement was **wrong** (a fresh world drawn per size compares universes, not sizes) and produced a curve that wandered 58/38/52/31.

Fund the test the portfolio learns most from. It is rarely the biggest brand.

The constraint nobody designs around: **you can only run an experiment on a brand that is already live.** Pre-launch brands carry the most risk and have no way to test. So "which live brand should we test" has an answer nobody computes — test the live brand that is the best lens on the launches queued behind it, even when that brand is small.

POLICY	PICKS	LEARNING BOUGHT	TO OTHERS
Biggest brand (what happens today)	brand 01	₹0.01L	—
Each P&L optimises its own	brand 03	₹0.24L	0%
Brand Zero	brand 00	₹1.87L	97%

Brand Zero's pick ranks **25th of 25 by revenue at stake** — literally the last brand anyone would think to test. It wins because the brands it teaches cannot yet be tested at all.

The honest half. Across 12 draws it picks a different test than "fund the biggest brand" in **5**. In the other 7, the flagship's own call happens to sit near its decision threshold and instinct is right. It earns its place on the draws where instinct is wrong, not by being different every time.

THE GOVERNANCE THAT MAKES IT SHIPPABLE

Telling Brand 7's lead "*your budget went to Brand 12 for portfolio value*" gets this switched off inside a quarter. They carry a P&L.

So the Allocator owns no brand budget. It owns a small ring-fenced learning budget and **bids** — offering to fund tests inside brands. Brands opt in because it is free money.

Same mathematics. Survivable politics.

allocator.py — expected value of information by Gauss–Hermite quadrature. The cross-brand covariance it needs falls straight out of the fit; no second model.

The one organ that pays at eleven brands. And at one.

Everything else needs neighbours before it works. This needs only that people write down what they expect *before* they find out — and get scored afterwards. The model forecasts too, on the same scale, scored the same way, with no special treatment.

FORECASTER	SAYS 80%	REALLY IS	RANGE WIDTH
Brand Zero	80%	82%	0.25
Lead A	80%	90%	0.40
Lead B	80%	5%	0.12

Lead B is not a worse operator. They state 12-point ranges on a question that moves 60 points. The score notices. That is a conversation about uncertainty, not about judgement — and it needs no sibling brands, which is why it ships in **week one**.

- **It is a proper scoring rule.** Quoting enormous ranges to be safe costs you on width; quoting tight ones costs you on misses. Honest reporting is the way to win.
- **Within two quarters you know who is calibrated on what.** Not who is smart — who is well-calibrated on price versus creative versus channel. Those are different people.
- **It turns judgement into evidence.** A calibrated operator's opinion can enter the ledger with a defensible weight instead of a guessed one.
- **It is also the honesty mechanism.** When the system refuses to answer and an operator overrides it — which is often correct — the override automatically opens a prediction in their name. Overriding is free. Overriding without saying what you expect is not.

Hold one brand out. Predict it cold. Repeat for all of them.

Remove every observation belonging to one brand, fit on the rest, predict it from nothing — exactly the position a brand is in on launch day. The number to watch is **coverage**: a range sold as 80% should contain the truth about 80% of the time.

AD MESSAGES — A DECISION THAT TRANSFERS

	ERROR	COVERAGE	SCORE
Each brand alone	0.53	90%	1.95
One average for all	0.53	5%	4.31
Brand Zero	0.27	70%	0.97

PRICE LADDERS — A DECISION THAT DOES NOT

	ERROR	COVERAGE	SCORE
Each brand alone	1.03	80%	3.86
One average for all	1.01	0%	7.96
Brand Zero	0.88	85%	3.07

The result we kept because it is unflattering. On price ladders — the class that does not transfer — Brand Zero is close to level with a brand that starts cold, and on some seeds slightly behind it. It has nothing to add there and does not pretend to. A system that won everywhere would be the less believable one.

WHAT THIS DOES NOT SHOW

These numbers are synthetic. The world they come from was generated with the same three-level structure the model fits. That demonstrates the mechanism; it is not evidence Think9's portfolio behaves this way.

The protocol that would settle it is **this exact backtest pointed at the real ledger** — a week-3 measurement, and it is in the plan. If it comes back weak, that is a finding to act on, not to explain away.

Where claims come from, and where humans decide.

SOURCES — ALL OF THEM ALREADY EXIST

- **Meta / Google Ads** — creative tests, the cleanest high-tier evidence
- **Shopify, Amazon, Flipkart** — price, pack, cohort revenue
- **Slack, meeting transcripts, vendor WhatsApp** — the reasoning no dashboard holds
- **Broadway & Brand Bridge** — 350+ third-party brands and 10k+ stores Think9 observes but does not own. **The answer to cold start:** the prior for brand 12 need not come only from eleven brands.

AGENTS — ONE JOB EACH

- **Extractors** — Claude with the claim schema as a forced output format, so a malformed claim is impossible. First job is to *reject*.
- **Adjudicator** — dedupes, resolves conflicts, assigns the evidence tier
- **Prior service** — nightly refit, per decision class
- **Query agent** — answers in Slack, cites its claims, refuses when it should
- **Allocator** — weekly ranked bid list. Never moves money.

HUMAN CHECKPOINTS — SIX

- **Claim confirmation** — brand lead, one click. Ignored → enters one tier *lower*.
- **Tier promotion** — only a named human can call something randomised. Tiers fall automatically; they never rise.
- **Refusal override** — allowed, logged, opens a prediction
- **Experiment funding** — weekly council. Nothing funded by inaction.
- **Pre-registration** — the one hard gate, above a spend threshold
- **Audit** — any number, back to its source, in two clicks

Each checkpoint specifies what happens by default if nobody acts — because in practice most checkpoints resolve by default, and a design that ignores that is decoration.

Stack: Postgres · Airbyte · Claude with structured outputs · numpy · Prefect · FastAPI + Slack. Deliberately no RAG and no vector store — the unit is a typed claim and the retrieval key is structured, so a semantic layer would be reaching for a familiar tool rather than the right one.

Sequenced by what pays at eleven brands — so pooling ships third.

WEEK 1

The ledger exists, and pre-registration goes live

One Postgres table. **20–30 past decisions backfilled by hand** by someone who was in the room — the most valuable day of the month, and it involves no AI. A Slack `/predict` command.

WEEK 2

Two extractors, and the first fit

Meta Ads and Shopify. Adjudicator plus the confirmation card. **Expect the first fit to be uninformative** — at this size wide intervals are the correct answer. Shipping it proves the pipeline, not the numbers.

WEEK 3

The honest measurement

Hold-one-brand-out against the **real** ledger. The query surface in Slack, refusal path included. This is the week the thesis gets tested.

WEEK 4

Allocation, and the calibration leaderboard

Ring-fenced budget agreed with finance; first council. One experiment funded that no single P&L would have paid for.

DELIBERATELY NOT BUILT IN 30 DAYS

WhatsApp and transcript extraction · quick-commerce ingestion · Broadway as a prior source · **anything autonomous** · cross-category transfer, because the model will almost certainly say it does not transfer and building for it before measuring would be exactly the mistake this system prevents.

A roadmap without exclusions is a wish.

HOW TO TELL WHETHER IT WORKED, AT DAY 30

- 1 · Is the ledger fed without being chased?
- 2 · Is anyone calibrated? Even one miscalibrated operator identified is worth the month.
- 3 · Does it beat starting cold on **real** claims? If not, say so.
- 4 · Did the Allocator fund something nobody else would have?

Every other asset in the portfolio depreciates. **Brand Zero appreciates with every launch.**

Inventory, creative, packaging tooling — even the brands themselves. Competitors can buy the same models tomorrow. They cannot buy eleven correlated brands' worth of pre-registered predictions and measured outcomes.

And the honest version of the promise: at eleven brands the ledger and the calibration layer carry the value, pooling compounds later, and the system's most useful habit is telling you when it has nothing to add.

Code & tests

github.com/SatyamSingh-Git/brand-zero
python demo.py · python cold_start.py

Interactive

sim/index.html — one file, no install
Four steps, about two minutes

Detail

[docs/architecture.md](https://docs.brand-zero.com/architecture.md)
[docs/roadmap-30-day.md](https://docs.brand-zero.com/roadmap-30-day.md)